

MongoDB Plugin & Toolchain for Hadoop

Brendan W. McAdams

10gen, Inc.

Dec. 3, 2010 @ MongoSV



Hadoop Explained...

- Started in February 2006 as part of the Apache Lucene project
- Based upon Google's MapReduce and GFS Papers
- Allows distributed, scalable data processing of huge datasets
- Java based, but with support for other JVM and Non-JVM Languages
- Lots of ecosystem tools to simplify interaction such as Pig and Hive
- In use at New York Times, Last.fm, Yahoo!, Amazon, Facebook and *many more companies*...
- Great tools for temporary Hadoop clusters such as the Cloudera Cluster Tools, Apache Whirr and Amazon's Elastic MapReduce.



Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- **Language: JavaScript only; not everyone wants to write JavaScript for data processing.**
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Why Integrate MongoDB?

- Language: JavaScript only; not everyone wants to write JavaScript for data processing.
 - Static Typing: JavaScript is dynamically typed which has limited value for jobs where a stronger type system is desired.
 - JVM Ecosystem: The JVM has a large number of libraries available for assisting in calculations and analysis which are available from Hadoop.
- Concurrency: Current JS Implementation is limited to one JS execution per server at a time.
- Scalability: Not a lot of ability to scale MongoDB's MapReduce except in cases of sharding.
- Integration: More development shops are integrating their data analysis into Hadoop jobs—opportunity to merge data from multiple systems including MongoDB for analysis.
- Hadoop is not a panacea—queries won't necessarily be fast(er) but it can handle larger scale and free up your MongoDB servers for long jobs.

Mongo + Hadoop ...

- Integrating MongoDB and Hadoop to read & write data from/to MongoDB via Hadoop
- 10gen has been working on a plugin to integrate the two systems, written in Pure Java
- About 6 months ago I explored the idea in Scala with a project called 'Luau'
- Support for pure MapReduce as well as Pig (Currently output only - input coming soon)
- With Hadoop Streaming (soon), write your MapReduce in Python, Ruby or Perl



MongoDB + MapReduce in Java I

We're going to walk through a sample application which parses US Treasury Bond historical Bid Curves since January 1990, and calculates an annual average for the 10 year Treasury.

- A sample of our dataset:

code/sample_treasury.js

```
0 { "_id" : { "$date" : 631238400000 }, "dayOfWeek" : "TUESDAY",  
  "bc3Year" : 7.9, "bc5Year" : 7.87, "bc10Year" : 7.94,  
  "bc20Year" : null, "bc1Month" : null, "bc2Year" : 7.87,  
  "bc3Month" : 7.83, "bc30Year" : 8, "bc1Year" : 7.81,  
  "bc7Year" : 7.98, "bc6Month" : 7.89 }  
1 { "_id" : { "$date" : 631324800000 }, "dayOfWeek" : "WEDNESDAY",  
  "bc3Year" : 7.96, "bc5Year" : 7.92, "bc10Year" : 7.99,  
  "bc20Year" : null, "bc1Month" : null, "bc2Year" : 7.94,  
  "bc3Month" : 7.89, "bc30Year" : 8.039999999999999,  
  "bc1Year" : 7.85, "bc7Year" : 8.039999999999999, "bc6Month"  
  : 7.94 }
```



MongoDB + MapReduce in Java II

```
2 { "_id" : { "$date" : 631411200000 }, "dayOfWeek" : "THURSDAY",  
  "bc3Year" : 7.93, "bc5Year" : 7.91, "bc10Year" : 7.98,  
  "bc20Year" : null, "bc1Month" : null, "bc2Year" : 7.92,  
  "bc3Month" : 7.84, "bc30Year" : 8.039999999999999,  
  "bc1Year" : 7.82, "bc7Year" : 8.02, "bc6Month" : 7.9 }  
3 { "_id" : { "$date" : 631497600000 }, "dayOfWeek" : "FRIDAY",  
  "bc3Year" : 7.94, "bc5Year" : 7.92, "bc10Year" : 7.99,  
  "bc20Year" : null, "bc1Month" : null, "bc2Year" : 7.9,  
  "bc3Month" : 7.79, "bc30Year" : 8.06, "bc1Year" : 7.79,  
  "bc7Year" : 8.029999999999999, "bc6Month" : 7.85 }
```



MongoDB + MapReduce in Java III

- The MongoDB JavaScript version of our demo:

code/mongo_treasury_mr.js

```
0 function m() {  
1     emit( this._id.getYear(), { count: 1, sum: this.bc10Year } )  
2 }  
3  
4 function r( year, values ) {  
5     var n = { count: 0, sum: 0 }  
6     for ( var i = 0; i < values.length; i++ ){  
7         n.sum += values[i].sum;  
8         n.count += values[i].count;  
9     }  
10  
11     return n;  
12 }  
13  
14 function f( year, value ){  
15     value.avg = value.sum / value.count;  
16     return value.avg;  
17 }
```



MongoDB + MapReduce in Java IV

```
18  
19 result = db.yield_historical.in.mapReduce( m , r, { finalize: f  
    });
```



MongoDB + MapReduce in Java V

- The same work can be done in Hadoop. . .



MongoDB Hadoop Job for Treasury Data I

Configuration

The Configuration is specified in XML

code/mongo-treasury_yield.xml

```
0 <?xml version="1.0"?>
1 <configuration>
2   <property>
3     <!-- run the job verbosely ? -->
4     <name>mongo.job.verbose</name>
5     <value>true</value>
6   </property>
7   <property>
8     <!-- Run the job in the foreground and wait for response, or
9     background it? -->
10    <name>mongo.job.background</name>
11    <value>false</value>
12  </property>
13  <property>
14    <!-- If you are reading from mongo, the URI -->
15    <name>mongo.input.uri</name>
16    <value>mongodb://localhost/demo.yield_historical.in</value>
```



MongoDB Hadoop Job for Treasury Data II

Configuration

```
16 </property>
17 <property>
18   <!-- If you are writing to mongo, the URI -->
19   <name>mongo.output.uri</name>
20   <value>mongodb://localhost/demo.yield_historical.out</value>
21 </property>
22 <property>
23   <!-- The query, in JSON, to execute [OPTIONAL] -->
24   <name>mongo.input.query</name>
25   <!--<value>{"x": {"$regex": "^eliot", "$options": ""}}</value>-->
26   <value></value>
27 </property>
28 <property>
29   <!-- The fields, in JSON, to read [OPTIONAL] -->
30   <name>mongo.input.fields</name>
31   <value></value>
32 </property>
33 <property>
34   <!-- A JSON sort specification for read [OPTIONAL] -->
35   <name>mongo.input.sort</name>
36   <value></value>
```



MongoDB Hadoop Job for Treasury Data III

Configuration

```
37 </property>
38 <property>
39   <!-- The number of documents to limit to for read [OPTIONAL] -->
40   <name>mongo.input.limit</name>
41   <value>0</value> <!-- 0 == no limit -->
42 </property>
43 <property>
44   <!-- The number of documents to skip in read [OPTIONAL] -->
45   <!-- TODO - Are we running limit() or skip() first? -->
46   <name>mongo.input.skip</name>
47   <value>0</value> <!-- 0 == no skip -->
48 </property>
49 <property>
50   <!-- Class for the mapper -->
51   <name>mongo.job.mapper</name>
52   <value>com.mongodb.hadoop.examples.TreasuryYieldMapper</value>
53 </property>
54 <property>
55   <!-- Reducer class -->
56   <name>mongo.job.reducer</name>
57   <value>com.mongodb.hadoop.examples.TreasuryYieldReducer</value>
```



MongoDB Hadoop Job for Treasury Data IV

Configuration

```
58 </property>
59 <property>
60   <!-- InputFormat Class -->
61   <name>mongo.job.input.format</name>
62   <value>com.mongodb.hadoop.MongoInputFormat</value>
63 </property>
64 <property>
65   <!-- OutputFormat Class -->
66   <name>mongo.job.output.format</name>
67   <value>com.mongodb.hadoop.MongoOutputFormat</value>
68 </property>
69 <property>
70   <!-- Output key class for the output format -->
71   <name>mongo.job.output.key</name>
72   <value>org.apache.hadoop.io.IntWritable</value>
73 </property>
74 <property>
75   <!-- Output value class for the output format -->
76   <name>mongo.job.output.value</name>
77   <value>org.apache.hadoop.io.DoubleWritable</value>
78 </property>
```



MongoDB Hadoop Job for Treasury Data V

Configuration

```
79 <property>
80   <!-- Output key class for the mapper [optional] -->
81   <name>mongo.job.mapper.output.key</name>
82   <value></value>
83   <value>org.apache.hadoop.io.IntWritable</value>
84 </property>
85 <property>
86   <!-- Output value class for the mapper [optional] -->
87   <name>mongo.job.mapper.output.value</name>
88   <value>org.apache.hadoop.io.DoubleWritable</value>
89 </property>
90 <property>
91   <!-- Class for the combiner [optional] -->
92   <name>mongo.job.combiner</name>
93   <value>com.mongodb.hadoop.examples.TreasuryYieldReducer</value>
94 </property>
95 <property>
96   <!-- Partitioner class [optional] -->
97   <name>mongo.job.partitionner</name>
98   <value></value>
99 </property>
```



MongoDB Hadoop Job for Treasury Data VI

Configuration

```
100 <property>
101   <!-- Sort Comparator class [optional] -->
102   <name>mongo.job.sort_comparator</name>
103   <value></value>
104 </property>
105
106 </configuration>
```



MongoDB Hadoop Job for Treasury Data I

Configuration

The Job is loaded & configured via a Java Tool:

code/TreasuryYieldXMLConfig.java

```
0 // TreasuryYieldXMLConfig.java
1 /*
2  * Copyright 2010 10gen Inc.
3  *
4  * Licensed under the Apache License, Version 2.0 (the "License");
5  * ...
6  */
7 package com.mongodb.hadoop.examples;
8
9 import java.io.*;
10 import java.util.*;
11
12 import org.apache.hadoop.conf.*;
13 import org.apache.hadoop.io.*;
14 import org.apache.hadoop.mapreduce.*;
15 import org.apache.hadoop.util.*;
16 import org.bson.*;
```



MongoDB Hadoop Job for Treasury Data II

Configuration

```
17
18 import com.mongodb.hadoop.util.*;
19
20 public class TreasuryYieldXMLConfig extends MongoTool {
21
22     static {
23         // Load the XML config defined in hadoop-local.xml
24         Configuration.addDefaultResource(
25             "src/examples/hadoop-local.xml" );
26         Configuration.addDefaultResource(
27             "src/examples/mongo-defaults.xml" );
28     }
29
30     public static void main( String[] args ) throws Exception{
31         final int exitCode = ToolRunner.run( new
32             TreasuryYieldXMLConfig(), args );
33         System.exit( exitCode );
34     }
35 }
```



MongoDB Hadoop Job for Treasury Data I

Configuration

MongoTool is a helper utility for doing XML configs with queries, etc:

code/MongoTool.java

```
0 // MongoTool.java
1 /*
2  * Copyright 2010 10gen Inc.
3  *
4  * Licensed under the Apache License, Version 2.0 (the "License");
5  * ...
6  */
7
8 package com.mongodb.hadoop.util;
9
10 import java.util.Map.Entry;
11
12 import org.apache.commons.logging.*;
13 import org.apache.hadoop.conf.*;
14 import org.apache.hadoop.mapreduce.*;
15 import org.apache.hadoop.util.*;
16
```



MongoDB Hadoop Job for Treasury Data II

Configuration

```
17 /**
18  * Tool for simplifying the setup and usage of Mongo Hadoop jobs
19  * using the Tool / Configured interfaces for use w/ a ToolRunner
20  * Primarily useful in cases of XML Config files.
21  *
22  * @author Brendan W. McAdams <brendan@10gen.com>
23  */
24 public class MongoTool extends Configured implements Tool {
25     private static final Log log = LoggerFactory.getLog( MongoTool.class
26     );
27
28     public int run( String[] args ) throws Exception{
29         /**
30          * ToolRunner will configure/process/setup the config
31          * so we need to grab the classlevel one
32          * This will be init'd with any loaded xml files or -D
33          prop=value params
34          */
35         final Configuration conf = getConf();
```



MongoDB Hadoop Job for Treasury Data III

Configuration

```
35     log.info( "Created a conf: '" + conf + "' on {" +
36     this.getClass() + "} as job named '" + _jobName + "'" );
37
38     for ( final Entry<String, String> entry : conf ) {
39         log.trace( String.format( "%s=%s\n", entry.getKey(),
40         entry.getValue() ) );
41     }
42
43     final Job job = new Job( conf , _jobName );
44     /**
45     * Any arguments specified with -D <property>=<value>
46     * on the CLI will be picked up and set here
47     * They override any XML level values
48     * Note that -D<space> is important - no space will
49     * not work as it get spicked up by Java itself
50     */
51     // TODO - Do we need to set job name somehow more
52     specifically?
53     // This may or may not be correct/sane
54     job.setJarByClass( this.getClass() );
55     final Class mapper = MongoConfigUtil.getMapper( conf );
```



MongoDB Hadoop Job for Treasury Data IV

Configuration

```
53     log.info( "Mapper Class: " + mapper );
54     job.setMapperClass( mapper );
55     job.setCombinerClass( MongoConfigUtil.getCombiner( conf ) );
56     job.setReducerClass( MongoConfigUtil.getReducer( conf ) );
57
58
59     job.setOutputFormatClass( MongoConfigUtil.getOutputFormat(
60     conf ) );
61     job.setOutputKeyClass( MongoConfigUtil.getOutputKey( conf )
62     );
63     job.setOutputValueClass( MongoConfigUtil.getOutputValue( conf
64     ) );
65
66     job.setInputFormatClass( MongoConfigUtil.getInputFormat( conf
67     ) );
68
69     /**
70      * Determines if the job will run verbosely e.g. print debug
71      * output
72      * Only works with foreground jobs
73      */
```



MongoDB Hadoop Job for Treasury Data V

Configuration

```
69         final boolean verbose = MongoConfigUtil.isJobVerbose( conf );
70         /**
71          * Run job in foreground aka wait for completion or
background?
72          */
73         final boolean background = MongoConfigUtil.isJobBackground(
conf );
74         try {
75             if ( background ) {
76                 log.info( "Setting up and running MapReduce job in
background." );
77                 job.submit();
78                 return 0;
79             }
80             else {
81                 log.info( "Setting up and running MapReduce job in
foreground, will wait for results. {Verbose? " + verbose + "}"
);
82                 return job.waitForCompletion( true ) ? 0 : 1;
83             }
84         }
```



MongoDB Hadoop Job for Treasury Data VI

Configuration

```
85         catch ( final Exception e ) {
86             log.error( "Exception while executing job... ", e );
87             return 1;
88         }
89     }
90
91     /**
92     * Main will be a necessary method to run the job - suggested
93     * implementation
94     * template:
95     * public static void main(String[] args) throws Exception {
96     * int exitCode = ToolRunner.run(new <YourClass>(), args);
97     * System.exit(exitCode);
98     * }
99     */
100
101     /**
102     * SET ME
103     * Defines the name of the job on the cluster.
104     * Left non-final to allow tweaking with serial #s, etc
```



MongoDB Hadoop Job for Treasury Data VII

Configuration

```
105     **/  
106     String _jobName = "<unnamed MongoTool job>";  
107 }
```



MongoDB Hadoop Job for Treasury Data I

Mapper

The Map function:

code/TreasuryYieldMapper.java

```
0 // TreasuryYieldMapper.java
1 /*
2  * Copyright 2010 10gen Inc.
3  *
4  * Licensed under the Apache License, Version 2.0 (the "License");
5  * ...
6  */
7 package com.mongodb.hadoop.examples;
8
9 import java.io.*;
10 import java.util.*;
11
12 import org.apache.commons.logging.*;
13
14 import org.apache.hadoop.conf.*;
15 import org.apache.hadoop.io.*;
16
```



MongoDB Hadoop Job for Treasury Data II

Mapper

```
17 import org.apache.hadoop.mapreduce.*;
18 import org.apache.hadoop.util.*;
19 import org.bson.*;
20
21 import com.mongodb.hadoop.util.*;
22
23 public class TreasuryYieldMapper extends Mapper<java.util.Date,
    BSONObject, IntWritable, DoubleWritable> {
24
25     private static final Log log = LogFactory.getLog(
    TreasuryYieldMapper.class );
26
27     public void map( java.util.Date key , BSONObject value , Context
    context ) throws IOException, InterruptedException{
28
29         int year = key.getYear() + 1900;
30         double bid10Year = ((Number) value.get( "bc10Year"
    )).doubleValue();
31
32         context.write( new IntWritable( year ), new DoubleWritable(
    bid10Year ) );
```



MongoDB Hadoop Job for Treasury Data III

Mapper

```
33  
34     }  
35  
36 }
```



MongoDB Hadoop Job for Treasury Data I

Reducer

The Reduce function:

code/TreasuryYieldReducer.java

```
0 // TreasuryYieldMapper.java
1 /*
2  * Copyright 2010 10gen Inc.
3  *
4  * Licensed under the Apache License, Version 2.0 (the "License");
5  * ...
6  */
7 package com.mongodb.hadoop.examples;
8
9 import java.io.*;
10 import java.util.*;
11
12 import org.apache.commons.logging.*;
13
14 import org.apache.hadoop.conf.*;
15 import org.apache.hadoop.io.*;
16
```



MongoDB Hadoop Job for Treasury Data II

Reducer

```
17 import org.apache.hadoop.mapreduce.*;
18 import org.apache.hadoop.util.*;
19 import org.bson.*;
20
21 import com.mongodb.hadoop.util.*;
22
23 public class TreasuryYieldReducer extends Reducer<IntWritable,
    DoubleWritable, IntWritable, DoubleWritable> {
24
25     private static final Log log = LogFactory.getLog(
    TreasuryYieldReducer.class );
26
27     public void reduce( IntWritable key , Iterable<DoubleWritable>
    values , Context context ) throws IOException,
    InterruptedException{
28
29         int count = 0;
30         double sum = 0;
31         for ( final DoubleWritable value : values ) {
32             log.debug( "Key: " + key + " Value: " + value );
33             sum += value.get();
```



MongoDB Hadoop Job for Treasury Data III

Reducer

```
34         count++;
35     }
36
37     double avg = sum / count;
38
39     log.info( "Average 10 Year Treasury for " + key.get() + " was
40 " + avg );
41
42     context.write( key, new DoubleWritable( avg ) );
43 }
44 }
```



MongoDB Hadoop Job For Treasury Data

Running

- Make sure the Java Driver is in the 'lib' directory on every cluster member, e.g. `sudo cp lib/mongo-java-driver-2.3.jar /usr/lib/hadoop-0.20/lib/` (Bounce hadoop after adding libraries)
- Build the Hadoop Jar (`ant resolve; ant jar`)
- Run the Job!



A Modified Pig Tutorial with MongoDB Output I

- We took the stock Pig Tutorial and modified it to save to MongoDB
- Input file is an anonymized Search Engine Log:

code/excite_log_snippet.txt

0	2A9EABFB35F5B954	970916105432	+md foods +proteins
1	BED75271605EBD0C	970916001949	yahoo chat
2	BED75271605EBD0C	970916001954	yahoo chat
3	BED75271605EBD0C	970916003523	yahoo chat
4	BED75271605EBD0C	970916011322	yahoo search
5	BED75271605EBD0C	970916011404	yahoo chat
6	BED75271605EBD0C	970916011422	yahoo chat
7	BED75271605EBD0C	970916012756	yahoo caht
8	BED75271605EBD0C	970916012816	yahoo chat
9	BED75271605EBD0C	970916023603	yahoo chat
10	BED75271605EBD0C	970916025458	yahoo caht
11	BED75271605EBD0C	970916025516	yahoo chat
12	BED75271605EBD0C	970916030348	yahoo chat
13	BED75271605EBD0C	970916034807	yahoo chat
14	BED75271605EBD0C	970916040755	yahoo chat
15	BED75271605EBD0C	970916090700	hawaii chat universe



A Modified Pig Tutorial with MongoDB Output II

16	BED75271605EBD0C	970916094445	yahoo chat
17	BED75271605EBD0C	970916191427	yahoo chat
18	BED75271605EBD0C	970916201045	yahoo chat
19	BED75271605EBD0C	970916201050	yahoo chat
20	BED75271605EBD0C	970916201927	yahoo chat
21	824F413FA37520BF	970916184809	garter belts
22	824F413FA37520BF	970916184818	garter belts
23	824F413FA37520BF	970916184939	lingerie
24	824F413FA37520BF	970916185051	spiderman

- The Script is only lightly modified from stock pig:


code/pigtutorial.pig

```
0 /*
1  * Licensed to the Apache Software Foundation (ASF) under one
2  * or more contributor license agreements. See the NOTICE file
3  * distributed with this work for additional information
4  * regarding copyright ownership. The ASF licenses this file
5  * to you under the Apache License, Version 2.0 (the
6  * "License"); you may not use this file except in compliance
```



A Modified Pig Tutorial with MongoDB Output III

```
7  * with the License.  You may obtain a copy of the License at
8  *
9  *      http://www.apache.org/licenses/LICENSE-2.0
10 *
11 * Unless required by applicable law or agreed to in writing,
12 * software
13 * distributed under the License is distributed on an "AS IS"
14 * BASIS,
15 * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express
16 * or implied.
17 * See the License for the specific language governing
18 * permissions and
19 * limitations under the License.
20 */
21
22 -- Query Phrase Popularity (local mode)
23
24 -- This script processes a search query log file from the Excite
25 search engine and finds search phrases that occur with
26 particular high frequency during certain times of the day.
```



A Modified Pig Tutorial with MongoDB Output IV

```
22 -- Register the tutorial JAR file so that the included UDFs can
    be called in the script.
23
24 -- Based on the Pig tutorial ,modified for Mongo support tests
25 REGISTER examples/pigtutorial/lib/pigtutorial.jar;
26 REGISTER mongo-hadoop.jar;
27 REGISTER lib/mongo-java-driver-2.3.jar;
28
29 -- Use the PigStorage function to load the excite log file into
    the raw bag as an array of records.
30 -- Input: (user,time,query)
31 raw = LOAD 'excite-small.log' USING PigStorage('\t') AS (user,
    time, query);
32
33 -- Call the NonURLDetector UDF to remove records if the query
    field is empty or a URL.
34 clean1 = FILTER raw BY
    org.apache.pig.tutorial.NonURLDetector(query);
35
36 -- Call the ToLower UDF to change the query field to lowercase.
37 clean2 = FOREACH clean1 GENERATE user, time,
    org.apache.pig.tutorial.ToLower(query) as query;
```



A Modified Pig Tutorial with MongoDB Output V

```
38
39 -- Because the log file only contains queries for a single day,
    we are only interested in the hour.
40 -- The excite query log timestamp format is YYMMDDHHMMSS.
41 -- Call the ExtractHour UDF to extract the hour (HH) from the
    time field.
42 houred = FOREACH clean2 GENERATE user,
    org.apache.pig.tutorial.ExtractHour(time) as hour, query;
43
44 -- Call the NGramGenerator UDF to compose the n-grams of the
    query.
45 ngramed1 = FOREACH houred GENERATE user, hour,
    flatten(org.apache.pig.tutorial.NGramGenerator(query)) as
    ngram;
46
47 -- Use the DISTINCT command to get the unique n-grams for all
    records.
48 ngramed2 = DISTINCT ngramed1;
49
50 -- Use the GROUP command to group records by n-gram and hour.
51 hour_frequency1 = GROUP ngramed2 BY (ngram, hour);
52
```



A Modified Pig Tutorial with MongoDB Output VI

```
53 -- Use the COUNT function to get the count (occurrences) of each
    n-gram.
54 hour_frequency2 = FOREACH hour_frequency1 GENERATE flatten($0),
    COUNT($1) as count;
55
56 -- Use the GROUP command to group records by n-gram only.
57 -- Each group now corresponds to a distinct n-gram and has the
    count for each hour.
58 uniq_frequency1 = GROUP hour_frequency2 BY group::ngram;
59
60 -- For each group, identify the hour in which this n-gram is
    used with a particularly high frequency.
61 -- Call the ScoreGenerator UDF to calculate a "popularity" score
    for the n-gram.
62 uniq_frequency2 = FOREACH uniq_frequency1 GENERATE flatten($0),
    flatten(org.apache.pig.tutorial.ScoreGenerator($1));
63
64 -- Use the FOREACH-GENERATE command to assign names to the
    fields.
65 uniq_frequency3 = FOREACH uniq_frequency2 GENERATE $1 as hour,
    $0 as ngram, $2 as score, $3 as count, $4 as mean;
66
```



A Modified Pig Tutorial with MongoDB Output VII

```
67 -- Use the FILTER command to move all records with a score less
    than or equal to 2.0.
68 filtered_uniq_frequency = FILTER uniq_frequency3 BY score > 2.0;
69
70 -- Use the ORDER command to sort the remaining records by hour
    and score.
71 ordered_uniq_frequency = ORDER filtered_uniq_frequency BY hour,
    score;
72
73 -- Use the PigStorage function to store the results.
74 -- Output: (hour, n-gram, score, count,
    average_counts_among_all_hours)
75 STORE ordered_uniq_frequency INTO
    'mongodb://localhost/demo.pig.output' USING
    com.mongodb.hadoop.pig.MongoStorage;
```



Mongo Hadoop Plugin

- Starting today, you can download and use the Hadoop Plugin:
`https://github.com/mongodb/mongo-hadoop`
- Still an Alpha—*not* production ready.
- Continuing Development. . .
 - Working on support for Hadoop Streaming, which allows Python, Ruby, and Perl to be used instead of Java (Waiting to migrate to .21 for Binary support)
 - Pig Input support.
 - Support for Sharding concurrency (Once one of you asks for it <hint-hint>)
 - Exploring other complimentary toolchains like Cascading.
 - Tell us how you are using Hadoop or how you want to integration MongoDB + Hadoop. . . So we can focus on the areas you, the users, want and need.



Questions?

- Twitter: **@rit** | mongodb: **@mongodb** | 10gen: **@10gen**
- email: **brendan@10gen.com**
- mongo-hadoop ... Available in a Pre-Production Alpha:
<https://github.com/mongodb/mongo-hadoop>
- Report Bugs, request features:
<https://github.com/mongodb/mongo-hadoop/issues>
- Pressing Questions?
 - IRC - freenode.net **#mongodb**
 - MongoDB Users List -
<http://groups.google.com/group/mongodb-user>
- 10gen is *hiring!* We need smart engineers in both NY and Bay Area: <http://10gen.com/jobs>
- Up Next: 15 Minute Break followed by Mathias Stearn on “MongoDB Internals: The Storage Engine”

