

Scala with MongoDB

Brendan W. McAdams

Novus Partners, Inc.

MongoDB NY Users Group
Aug. 10, 2010



Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

What Is Scala?

- Stands for ‘scalable language’
- Blends functional & object oriented programming concepts
- Designed by Martin Odersky - author of the Generics in Java 5
- Compiles to Java Bytecode, allowing it to run on the JVM(some support for .Net CLR)
- At runtime it's just JVM Bytecode - can call Java objects and vice versa
- Alternate concurrency model based on Actors (as used in Erlang)
- Supports programming concepts like immutability and closures natively.
- Type inference allows reduction of unnecessary type annotations.
- Immutability built-in...
 - Declaring a variable `var` makes it mutable; `val` makes it immutable.
 - Two separated collection libraries -
`scala.collection.mutable` and
`scala.collection.immutable`

The Truth about Immutability vs. Mutability in Scala

... Most Scala developers learn to favor immutability.

Outline

- 1 Introduction
 - What is Scala?
 - **Java <-> Scala Basics**
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Object Oriented Programming in Scala

Three core object oriented declarations in Scala

- Classes
- Traits
- Objects

Object Oriented Programming in Scala

Three core object oriented declarations in Scala

- Classes
- Traits
- Objects

Object Oriented Programming in Scala

Three core object oriented declarations in Scala

- Classes
- Traits
- Objects

Object Oriented Programming in Scala

Three core object oriented declarations in Scala

- Classes
- Traits
- Objects

Classes in Scala I

- Classes are similar to most languages.
- Made up of properties and methods.
- An instance of a class is created with the “new” keyword.
- No multiple inheritance, but can implement traits.
- Constructors are part of the class definition, rather than separate methods.
- Overloaded constructors are allowed, but must call another constructor *before any other code is executed*.

Traits in Scala I

- Easiest way to think of a Trait is an interface with concrete declarations; allows abstract and concrete definitions.
- Traits can **NOT** be directly instantiated.
- Traits are stackable - multiple traits can call one another for specific functionality.
- Traits can be used as “mixins” including being added at class instantiation to change functionality.

Scala's Object Type I

- Scala does not support static method declaration in classes; any method declared in an “object” is static.
- Scala objects act as “companions” - if an object and class with the same name exist, they complement each other.
- Scala objects are a system managed singleton - only one instance of the object ever exists (From Java's view an “object” is actually <Name>\$)
- Companion objects are often used as factories to prevent/proxy direct class instantiation.

User Defined Operators I

- Scala allows the definition of operators...
- Not technically operator overloading, as the JVM doesn't have operators - they're language built-ins in Java, etc.
- In Scala, there are no built-in operators. Some are predefined for sanity (like `+`, `-`, `/` and `*` on Numeric types) but operators are just methods.
- Scala allows any operator to be defined by the user - including some special ones like unaries (`+foo`, `-foo`).
- Syntactic Sugar: To facilitate statements like `foo + bar` Scala allows methods to be called without the `.` or parentheses. . . Useful for DSLs and fluid syntaxes!

Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - **Implicits and Pimp Hats**
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Functional Programming and Scala

- What is Functional Programming?
 - Functions are Objects
 - Immutability
- A few crucial Scala concepts which depend upon FP (and Scala programmers delight in)
 - Anonymous Functions
 - `apply()` (and `unapply`)
 - Useful tools like `group`, `foreach` and `map`

Outline

- 1 Introduction
 - What is Scala?
 - Java <-> Scala Basics
 - Implicits and Pimp Hats
 - **What is MongoDB?**
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Helping Java + Scala Interact

- Implicits, “Pimp My Library” and various conversion helper tools simplify the work of interacting with Java.
- Scala and Java have their own completely different collection libraries.
- Some builtins ship with Scala to make this easier.

Helping Java + Scala Interact

- Implicits, “Pimp My Library” and various conversion helper tools simplify the work of interacting with Java.
- Scala and Java have their own completely different collection libraries.
- Some builtins ship with Scala to make this easier.

Helping Java + Scala Interact

- Implicits, “Pimp My Library” and various conversion helper tools simplify the work of interacting with Java.
- Scala and Java have their own completely different collection libraries.
- Some builtins ship with Scala to make this easier.

Interoperability in Scala 2.7.x

- **Scala 2.7.x shipped with `scala.collection.jcl`.**
- `scala.collection.jcl.Conversions` contained some implicit converters, but only to and from the wrapper versions - no support for “real” Scala collections.
- Neglected useful base interfaces like `Iterator` and `Iterable`
- @jorgeortiz85 provided `scala-javautils`, which used “Pimp My Library” to do a better job.

Interoperability in Scala 2.7.x

- Scala 2.7.x shipped with `scala.collection.jcl`.
- `scala.collection.jcl.Conversions` contained some implicit converters, but only to and from the wrapper versions - no support for “real” Scala collections.
- Neglected useful base interfaces like `Iterator` and `Iterable`
- @jorgeortiz85 provided `scala-javautils`, which used “Pimp My Library” to do a better job.

Interoperability in Scala 2.7.x

- Scala 2.7.x shipped with `scala.collection.jcl`.
- `scala.collection.jcl.Conversions` contained some implicit converters, but only to and from the wrapper versions - no support for “real” Scala collections.
- Neglected useful base interfaces like `Iterator` and `Iterable`
- @jorgeortiz85 provided `scala-javautils`, which used “Pimp My Library” to do a better job.

Interoperability in Scala 2.8.x

- Scala 2.8.x improves the interop game significantly.
- JCL is gone - focus has shifted to proper interoperability w/ built-in types.
- `scala.collection.jcl.Conversions` replaced by `scala.collection.JavaConversions` - provides implicit conversions to & from Scala & Java Collections.
- Includes support for the things missing in 2.7 (`Iterable`, `Iterator`, etc.)
- Great for places where the compiler can guess what you want (implicits); falls short in some cases (like BSON Encoding, as we found in Casbah)
- @jorgeortiz85 has updated `scala-javautils` for 2.8 with `scalaj-collection`
- Explicit `asJava` / `asScala` methods for conversions. Adds `foreach` method to Java collections.

Interoperability in Scala 2.8.x

- Scala 2.8.x improves the interop game significantly.
- JCL is gone - focus has shifted to proper interoperability w/ built-in types.
- `scala.collection.jcl.Conversions` replaced by `scala.collection.JavaConversions` - provides implicit conversions to & from Scala & Java Collections.
- Includes support for the things missing in 2.7 (`Iterable`, `Iterator`, etc.)
- Great for places where the compiler can guess what you want (implicits); falls short in some cases (like BSON Encoding, as we found in Casbah)
- @jorgeortiz85 has updated `scala-javautils` for 2.8 with `scalaj-collection`
- Explicit `asJava` / `asScala` methods for conversions. Adds `foreach` method to Java collections.

Interoperability in Scala 2.8.x

- Scala 2.8.x improves the interop game significantly.
- JCL is gone - focus has shifted to proper interoperability w/ built-in types.
- `scala.collection.jcl.Conversions` replaced by `scala.collection.JavaConversions` - provides implicit conversions to & from Scala & Java Collections.
- Includes support for the things missing in 2.7 (`Iterable`, `Iterator`, etc.)
- Great for places where the compiler can guess what you want (implicits); falls short in some cases (like BSON Encoding, as we found in Casbah)
- @jorgeortiz85 has updated `scala-javautils` for 2.8 with `scalaj-collection`
- Explicit `asJava` / `asScala` methods for conversions. Adds `foreach` method to Java collections.

Interoperability in Scala 2.8.x

- Scala 2.8.x improves the interop game significantly.
- JCL is gone - focus has shifted to proper interoperability w/ built-in types.
- `scala.collection.jcl.Conversions` replaced by `scala.collection.JavaConversions` - provides implicit conversions to & from Scala & Java Collections.
- Includes support for the things missing in 2.7 (`Iterable`, `Iterator`, etc.)
- Great for places where the compiler can guess what you want (implicits); falls short in some cases (like BSON Encoding, as we found in Casbah)
- @jorgeortiz85 has updated `scala-javautils` for 2.8 with `scalaj-collection`
- Explicit `asJava` / `asScala` methods for conversions. Adds `foreach` method to Java collections.

Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - **A Taste of MongoDB**
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

So WTF is an 'Implicit', anyway?

- Implicit Arguments

- 'Explicit' arguments indicates a method argument you pass, well *explicitly*.
- 'Implicit' indicates a method argument which is... *implied*. (But you can pass them explicitly too.)
- Implicit arguments are passed in Scala as an additional argument list:

```
import com.mongodb._
import org.bson.types.ObjectId

def query(id: ObjectId)(implicit coll: DBCollection) = coll.findOne(id)

val conn = new Mongo()
val db = conn.getDB("test")
implicit val coll = db.getCollection("testData")

// coll is passed implicitly
query(new ObjectId())

// or we can override the argument
query(new ObjectId())(db.getCollection("testDataExplicit"))
```

- How does this differ from default arguments?

So WTF is an 'Implicit', anyway?

- Implicit Arguments

- 'Explicit' arguments indicates a method argument you pass, well *explicitly*.
- 'Implicit' indicates a method argument which is... *implied*. (But you can pass them explicitly too.)
- Implicit arguments are passed in Scala as an additional argument list:

```
import com.mongodb._
import org.bson.types.ObjectId

def query(id: ObjectId)(implicit coll: DBCollection) = coll.findOne(id)

val conn = new Mongo()
val db = conn.getDB("test")
implicit val coll = db.getCollection("testData")

// coll is passed implicitly
query(new ObjectId())

// or we can override the argument
query(new ObjectId())(db.getCollection("testDataExplicit"))
```

- How does this differ from default arguments?

So WTF is an ‘Implicit’, anyway?

- Implicit Methods/Conversions

- If you try passing a type to a Scala method argument which doesn't match. . .

```
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212") // won't compile
```

- A fast and loose example, but simple. Fails to compile.
- But with implicit methods, we can provide a conversion path. . .

```
implicit def strToNum(x: String) = x.toInt
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212")
```

- In a dynamic language, this may be called “monkey patching”.
Unlike Perl, Python, etc. Scala resolves implicits at compile time.

So WTF is an ‘Implicit’, anyway?

- Implicit Methods/Conversions

- If you try passing a type to a Scala method argument which doesn't match. . .

```
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212") // won't compile
```

- A fast and loose example, but simple. Fails to compile.
- But with implicit methods, we can provide a conversion path. . .

```
implicit def strToNum(x: String) = x.toInt
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212")
```

- In a dynamic language, this may be called “monkey patching”.
Unlike Perl, Python, etc. Scala resolves implicits at compile time.

So WTF is an ‘Implicit’, anyway?

- Implicit Methods/Conversions

- If you try passing a type to a Scala method argument which doesn't match. . .

```
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212") // won't compile
```

- A fast and loose example, but simple. Fails to compile.
- But with implicit methods, we can provide a conversion path. . .

```
implicit def strToNum(x: String) = x.toInt
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212")
```

- In a dynamic language, this may be called “monkey patching”.
Unlike Perl, Python, etc. Scala resolves implicits at compile time.

So WTF is an ‘Implicit’, anyway?

- Implicit Methods/Conversions

- If you try passing a type to a Scala method argument which doesn't match. . .

```
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212") // won't compile
```

- A fast and loose example, but simple. Fails to compile.
- But with implicit methods, we can provide a conversion path. . .

```
implicit def strToNum(x: String) = x.toInt
def printNumber(x: Int) = println(x)

printNumber(5)
printNumber("212")
```

- In a dynamic language, this may be called “monkey patching”. Unlike Perl, Python, etc. Scala resolves implicits at compile time.

Pimp My Library I

- Coined by Martin Odersky in a 2006 Blog post. Similar to C# extension methods, Ruby modules.
- Uses implicit conversions to tack on new methods at runtime.
- Either return a new “Rich_” or anonymous class...

```
import com.mongodb.gridfs.{GridFSInputFile => MongoDBGridFSInputFile}

class GridFSInputFile protected[mongodb] (override val underlying:
  MongoDBGridFSInputFile) extends GridFSFile {
  def filename_=(name: String) = underlying.setFilename(name)
  def contentType_=(cT: String) = underlying.setContentType(cT)
}

object PimpMyMongo {
  implicit def mongoConnAsScala(conn: Mongo) = new {
    def asScala = new MongoClient(conn)
  }

  implicit def enrichGridFSInput(in: MongoDBGridFSInputFile) =
    new GridFSInputFile(in)
}

import PimpMyMongo._
```



Pimp My Library II

Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - **MongoDB + Scala Drivers**

- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka

- 3 Closing

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Introducing MongoDB

- Categorized as a “Document-Oriented Database”
 - Features of both Key-Value Stores & RDBMS’
 - Rich query interface.
 - Works with JSON-like Documents
 - Favors embedding related data over “foreign key” relationships
- Free license (A-GPL) cross-platform (Packages for Linux, Windows, Mac OS X, Windows, FreeBSD & Solaris)
- Cursor-based query results
- ServerSide Javascript
 - Stored Javascript functions server-side
 - Powerful aggregation - Map/Reduce, Group Commands
 - JS Statements in queries (no indexes though)
- Indexing system is much like RDBMS, includes Geospatial support.
- Scalable file storage with GridFS
- Data scalability with Replica Sets & Autosharding

Programming with MongoDB

- Provides a native API which allows interaction to adapt to the programming language (rather than vice versa).
- Official drivers for...
 - C
 - C++
 - Java
 - JavaScript
 - Perl
 - PHP
 - Python
 - Ruby
- Community supported drivers include...
 - .Net: C# & F#
 - JVM: Clojure, Scala, Groovy
 - Erlang
 - Haskell
 - Go
 - ...and many more.

Programming with MongoDB

- Provides a native API which allows interaction to adapt to the programming language (rather than vice versa).
- Official drivers for...
 - C
 - C++
 - Java
 - JavaScript
 - Perl
 - PHP
 - Python
 - Ruby
- Community supported drivers include...
 - .Net: C# & F#
 - JVM: Clojure, Scala, Groovy
 - Erlang
 - Haskell
 - Go
 - ...and many more.

Programming with MongoDB

- Provides a native API which allows interaction to adapt to the programming language (rather than vice versa).
- Official drivers for...
 - C
 - C++
 - Java
 - JavaScript
 - Perl
 - PHP
 - Python
 - Ruby
- Community supported drivers include...
 - .Net: C# & F#
 - JVM: Clojure, Scala, Groovy
 - Erlang
 - Haskell
 - Go
 - ... and many more.

But is anyone actually *using* it?!?

MongoDB is deployed in production at companies including...

- New York Times
- Foursquare
- bit.ly
- SourceForge
- Etsy
- Disqus
- Github
- ... The Large Hadron Collider.

But is anyone actually *using* it?!?

MongoDB is deployed in production at companies including...

- New York Times
- Foursquare
- bit.ly
- SourceForge
- Etsy
- Disqus
- Github
- ... The Large Hadron Collider.

Outline

- 1 Introduction
 - What is Scala?
 - Java \leftrightarrow Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Core Concepts

- MongoDB's equivalent to "tables" are called "collections"; "collections" contain "documents" (individual pieces of data)
- Databases & Collections are lazy - they are created when first inserted into.
- MongoDB's wire format/internal document representation is **BSON**. . .
 - **BSON** is a binary optimized flavor of **JSON**; corrects **JSON**'s inefficiency in string encoding (Base64).
 - Supports extras including Regular Expressions, Byte Arrays, DateTimes & Timestamps, as well as datatypes for Javascript code blocks & functions.
 - **BSON** implementation being split into its own package in most drivers.
 - Creative Commons licensed **<http://bsonspec.org>**
- Java driver represents **BSON** with a map-like **DBObject** (Which most Scala drivers use); many dynamic languages (Perl, Ruby, Python, etc) use native dictionary objects.

The basics of Querying I

- Find a single row with *findOne()*; returns the first document found (by natural order).
- You can find all documents matching your query with *find()*. No query means you get the entire collection back.
- Queries are specified as **BSON** documents to match against.
- The *find()* and *findOne()* methods can take an optional second **DBObject** specifying the fields to return.
- If you have an embedded object (for example, an address object) you can retrieve it with dot notation in the fields list (e.g. “*address.city*” retrieves just the city value).
- Use *limit()*, *skip()* and *sort()* on result objects (**DBCursor** in Java-driver land) to adjust your results. These all return a new cursor.

The basics of Querying II

- *distinct()* can be used (on **DBCollection** to find all distinct values for a given key; it returns a list of values.

```
> db.routes.findOne({"route_short_name": "E"})
{
  "_id" : ObjectId("4c5f755608c3693f59580f8c"),
  "route_id" : "E",
  "agency_id" : "MTA NYCT",
  "route_short_name" : "E",
  "route_long_name" : "8 Avenue Local",
  "route_desc" : "Trains operate between Jamaica Center (Parsons/Archer),
Queens, and World Trade Center, Manhattan, at all times.",
  "route_type" : 1,
  "route_url" : "http://www.mta.info/nyct/service/pdf/tecur.pdf"
}

> db.routes.find({"route_long_name": /Local$/},
  {"route_short_name": 1, "route_long_name": 1})

{ "_id" : ObjectId("4c5f755608c3693f59580f7f"), "route_short_name" : 1,
  "route_long_name" : "Broadway - 7 Avenue Local" }
{ "_id" : ObjectId("4c5f755608c3693f59580f84"), "route_short_name" : 6,
  "route_long_name" : "Lexington Avenue Local" }
{ "_id" : ObjectId("4c5f755608c3693f59580f86"), "route_short_name" : 7,
  "route_long_name" : "Flushing Local" }
{ "_id" : ObjectId("4c5f755608c3693f59580f8a"), "route_short_name" : "C",
  "route_long_name" : "8 Avenue Local" }
```

The basics of Querying III

```
{ "_id" : ObjectId("4c5f755608c3693f59580f8c"), "route_short_name" : "E",  
  "route_long_name" : "8 Avenue Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f8d"), "route_short_name" : "F",  
  "route_long_name" : "Queens Blvd Express/ 6 Av Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f91"), "route_short_name" : "J",  
  "route_long_name" : "Nassau St Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f92"), "route_short_name" : "L",  
  "route_long_name" : "14 St-Canarsie Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f93"), "route_short_name" : "M",  
  "route_long_name" : "Nassau St Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f96"), "route_short_name" : "R",  
  "route_long_name" : "Broadway Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f99"), "route_short_name" : "V",  
  "route_long_name" : "Queens Blvd/6 Av Local" }  
{ "_id" : ObjectId("4c5f755608c3693f59580f9a"), "route_short_name" : "W",  
  "route_long_name" : "Broadway Local" }  
  
> db.routes.distinct("route_short_name")  
[  
  1,  
  2,  
  3,  
  4,  
  5,  
  6,  
  7,  
  "A",  
  "B",  
  "C",
```

The basics of Querying IV

```
"D",  
"E",  
"F",  
"G",  
"J",  
/*...*/  
]
```


Query Operators I

- MongoDB is no mere Key-Value store. There are myriad powerful operators to enhance your MongoDB queries. . .
 - Conditional Operators: **\$gt** (>), **\$lt** (<), **\$gte** (>=), **\$lte** (<=)
 - Negative Equality: **\$ne** (!=)
 - Array Operators: **\$in** (SQL “IN” clause...takes an array), **\$nin** (Opposite of “IN”), **\$all** (Requires all values in the array match), **\$size** (Match the size of an array)
 - Field Defined: **\$exists** (boolean argument)(Great in a schemaless world)
 - Regular Expressions (Language dependent - most drivers support it)
 - Pass Arbitrary Javascript with **\$where** (No OR statements, so use WHERE for complex range filters)
 - Negate any operator with **\$not**
- Using a query operator requires nested objects. . .

Query Operators II

```
> db.stops.find({"stop_lat" : {$lt: 40.6}, {"stop_lon": {$gte: -73.8}}})
{ "_id" : ObjectId("4c5f755608c3693f59580ef0"), "stop_lat" : 40.590927, "stop_lon" :
  -73.796924, "stop_id" : "H06", "stop_name" : "BEACH 67TH ST - GASTON",
  "location_type" : 0, "stop_geo" : { "lat" : 40.590927, "lon" : -73.796924 } }
{ "_id" : ObjectId("4c5f755608c3693f59580ef1"), "stop_lat" : 40.592374, "stop_lon" :
  -73.788522, "stop_id" : "H07", "stop_name" : "BEACH 60TH ST - STRAITON AV",
  "location_type" : 0, "stop_geo" : { "lat" : 40.592374, "lon" : -73.788522 } }
{ "_id" : ObjectId("4c5f755608c3693f59580ef2"), "stop_lat" : 40.592943, "stop_lon" :
  -73.776013, "stop_id" : "H08", "stop_name" : "BEACH 44TH ST - FRANK AV",
  "location_type" : 0, "stop_geo" : { "lat" : 40.592943, "lon" : -73.776013 } }
{ "_id" : ObjectId("4c5f755608c3693f59580ef3"), "stop_lat" : 40.595398, "stop_lon" :
  -73.768175, "stop_id" : "H09", "stop_name" : "BEACH 36TH ST - EDGEMERE",
  "location_type" : 0, "stop_geo" : { "lat" : 40.595398, "lon" : -73.768175 } }

> db.trips.findOne({"route_id": {$in: ["E", "4", "5"]}})
{
  "_id" : ObjectId("4c5f755708c3693f59583400"),
  "route_id" : "E",
  "service_id" : "B20100308W",
  "trip_id" : "B20100308W_001350_E..S04R",
  "trip_headsign" : "To World Trade Ctr",
  "direction_id" : 1,
  "shape_id" : 177710
}

> db.trips.find({"route_id": {$in: ["E", "4", "5"]}}).count()
928
```

Query Operators III

- No syntactic sugar in Java to make it easier...

Insert/Update/Save I

- Objects in MongoDB Collections have an “_id” field, which must be unique.
- Three ways to add/update data in MongoDB. . .
 - `insert()` always attempts to add a new row. If “_id” is present and contains a value already in the collection, insert fails.
 - `save()` inserts if there is no “_id” field, otherwise it tries to update the document with the specified “_id”.
 - `update()` takes a query and the new values to save. By default it updates only the first document matching the query.
 - For `update()` you can specify two booleans whose default is false: *upsert*, which indicates you wish to create a new document if the query doesn't match, and *multi*, which allows updating **all** documents who match the query.

Insert/Update/Save II

```
> db.testData.insert({"userCount": 5})
> x = db.testData.findOne({"userCount": 5})
{ "_id" : ObjectId("4c607f48150c335a4e187f41"), "userCount" : 5 }
> x.userCount
5
> x.userCount = 20
20
> db.testData.save(x)
> db.testData.findOne({_id: x._id})
{ "_id" : ObjectId("4c607f48150c335a4e187f41"), "userCount" : 20 }
> db.testData.update({_id: x._id}, {$inc: {"userCount": 12}})
> db.testData.findOne({_id: x._id})
{ "_id" : ObjectId("4c607f48150c335a4e187f41"), "userCount" : 32 }
// upsert
> db.testData.update({"userCount": 5}, {"userCount": 209}, true)
> db.testData.findOne({"userCount": 209} )
{ "_id" : ObjectId("4c60800e08c3693f5962dda5"), "userCount" : 209 }
```

Geospatial Support I

- MongoDB supports Geospatial indexing and distance based queries
- I loaded all of the NYC Subway data (in Google Transit format) into MongoDB
- Quick python code to index the “Stops” data.

```
connection = Connection()
db = connection['nyct_subway']
print "Indexing the Stops Data."
for row in db.stops.find():
    row['stop_geo'] = {'lat': row['stop_lat'], 'lon': row['stop_lon']}
    db.stops.save(row)

db.stops.ensure_index([('stop_geo', pymongo.GEO2D)])
```

- “stop_geo” field is now Geospatially indexed for each stop.
- How hard is it to find the 5 closest subway stops to Meetup HQ?

Geospatial Support I

```
db.stops.find( { stop_geo: { $near: [40.726021, -73.99617] } }, {'stop_name':  
  1}).limit(5);  
  
  { "_id" : ObjectId("4c5f755608c3693f59580e9b"), "stop_name" : "BROADWAY-LAFAYETTE  
    " }  
  { "_id" : ObjectId("4c5f755608c3693f59580e29"), "stop_name" : "BLEECKER  
    STREET-LEXINGTON" }  
  { "_id" : ObjectId("4c5f755608c3693f59580f50"), "stop_name" : "PRINCE STREET  
    " }  
  { "_id" : ObjectId("4c5f755608c3693f59580e2a"), "stop_name" : "SPRING  
    STREET-LEXINGTON" }  
  { "_id" : ObjectId("4c5f755608c3693f59580f4f"), "stop_name" : "8TH STREET (NYU)      " }  
  }
```

- Further commands exist to define a rectangle or circle radius for the search.

Finally, Data Scalability.

- Traditional master-slave replication
- Replica Sets (new in 1.6)
 - Replaces master-slave setup with 1-7 server clusters
 - Automatic failover and recovery
- AutoSharding (new in 1.6)
 - Horizontal scaling - partition your collections & data across as many nodes as necessary.
 - Multiple nodes can service the same shard, allowing for balancing & failover.
 - Map/Reduce runs across multiple shards, allowing concurrency.

Outline

- 1 Introduction
 - What is Scala?
 - Java <-> Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers
- 2 Scala + MongoDB == Win
 - lift-mongo
 - casbah
 - STM + MongoDB via Akka
- 3 Closing

Using Scala with the official Java Driver I

- JVM Object are JVM Objects. . .

```
import com.mongodb._

val conn = new Mongo()
val db = conn.getDB("test")
val coll = db.getCollection("testData")

val pies = new BasicDBList()
pies.add("cherry")
pies.add("blueberry")
pies.add("apple")
pies.add("rhubarb")
pies.add("3.14")

val doc = new BasicDBObject()
doc.put("foo", "bar")
doc.put("spam", "eggs")
doc.put("up", "down")
doc.put("pie", pies)

coll.insert(doc)
```

- . . . Not terribly “Scala-ey”.

Using Scala with the official Java Driver II

- The Java driver works, but doesn't fit well in Scala.
- You need to convert your Scala objects to Java Objects, and get nothing but Java Objects out.
- Gets messy quickly.

The Scala Community Adapted... I

Compare the previous with various Scala drivers.

- mongo-scala-driver wraps & enhances the Java driver:

```
import com.mongodb._
import com.osinka.mongodb._

val conn = new Mongo()
val db = conn.getDB("test")
val coll = db.getCollection("testData").asScala

coll << Map(
  "foo" -> "bar",
  "spam" -> "eggs",
  "up" -> "down",
  "pie" -> List(
    "cherry",
    "blueberry",
    "apple",
    "rhubarb",
    "3.14"
  )
)
```

The Scala Community Adapted... II

- .. Much better, although I was confused initially. Has a object<->MongoDB mapping layer.
- lift-mongodb has more than one way to do it... here's just a taste:

```
import com.mongodb._

import net.liftweb.mongodb._
import net.liftweb.json._
import net.liftweb.json.JsonAST.JObject
import net.liftweb.json.JsonDSL._

implicit val formats = DefaultFormats.lossless

MongoDB.defineDb(DefaultMongoIdentifier,
  MongoAddress(MongoHost("localhost", 27017)), "test")

val json = JsonParser.parse("""
{ "foo": "bar",
  "spam": "eggs",
  "up": "down",
  "pie": [
    "cherry",
    "blueberry",
    "apple",
    "rhubarb",
    "3.14"
  ]
}
```

The Scala Community Adapted... III

```
    }  
  }  
  """).asInstanceOf[JObject]  
  
  MongoDB.useCollection("testData") { coll => {  
    coll.save(JObjectParser.parse(json))  
  }}  
}
```

- ... Lift's JS & JSON tools make it very flexible, as we'll see later. Also has an ActiveRecord style Object<->MongoDB Mapping layer.
- Casbah reflects my own attempt at creating a sane interface between Scala & MongoDB. Influenced by pymongo:

The Scala Community Adapted... IV

```
import com.novus.casbah.mongodb.Imports._

val coll = MongoConnection()("test")("testData")

val builder = MongoDBObject.newBuilder
builder += "foo" -> "bar"
builder += "spam" -> "eggs"
builder += "up" -> "down"
builder += "pie" -> List("cherry", "blueberry",
                        "apple", "rhubarb", "3.14")

coll += builder.result
```

- ... The syntax is still growing but is meant to match Scala syntax sanely. Object<->MongoDB Mapping coming soon.
- We're going to cover several tools, although I know Casbah best.

Outline

- 1 Introduction
 - What is Scala?
 - Java <-> Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers

- 2 **Scala + MongoDB == Win**
 - **lift-mongo**
 - casbah
 - STM + MongoDB via Akka

- 3 Closing

lift-mongo I

- Formerly “scamongo”, integrated with Lift as of 2.0
- Base code provides session wrappers to MongoDB, still utilizes Java driver'sDBObject code.

```
MongoDB.defineDb(DefaultMongoIdentifier,
                  MongoAddress(MongoHost("localhost", 27017)), "test")

MongoDB.useCollection(collectionName) ( coll => {
  val doc = new BasicDBObject
  doc.put("name", "MongoDB")
  doc.put("type", "database")
  doc.put("count", 1)
  // save the doc to the db
  coll.save(doc)
})

// Alternately, do everything in a single thread...
MongoDB.useSession ( db => {
  val coll = db.getCollection("testCollection")
  val doc = new BasicDBObject
  doc.put("name", "MongoSession")
  doc.put("type", "db")
  doc.put("count", 1)
  coll.save(doc)
})
```



lift-mongo II

- “lift-mongo-record” provides object mapping.
- No native query syntax, but Foursquare is working on open sourcing something they use internally.

lift-mongo-record & querying I

- Object definitions are fairly straightforward...

```
class MainDoc extends MongoRecord[MainDoc] with MongoId[MainDoc] {  
  def meta = MainDoc  
  object name extends StringField(this, 12)  
  object cnt extends IntField(this)  
  object refdoc extends DBRefField[MainDoc, RefDoc](this, RefDoc)  
  object refdocId extends ObjectIdField(this) {  
    def fetch = RefDoc.find(value)  
  }  
}  
  
object MainDoc extends MainDoc with MongoMetaRecord[MainDoc] {  
  def createRecord = new MainDoc  
}  
  
class RefDoc extends MongoRecord[RefDoc] with MongoId[RefDoc] {  
  def meta = RefDoc  
}  
  
object RefDoc extends RefDoc with MongoMetaRecord[RefDoc] {  
  def createRecord = new RefDoc  
}  
  
// Querying appears limited to constructing Mongo DBObjects  
val mdq1 = MainDoc.findAll(("name" -> "md1"))
```

lift-mongo-record & querying II

- Foursquare's query library allow for a saner way to query data. . .

```
// FSMongoRecord extends "MongoRecord" to add a few methods
class Venue extends FSMongoRecord[Venue] {
  def meta = Venue

  object venueName extends FSStringField(this, 255)
  object keywords extends StringField(this, 255)
  object userid extends LongField(this)
  object closed extends BooleanField(this) with AuditableField[Venue]
  object mayor LegacyForeignKey(this, User) {
    override def optional_? = true
  }
  object mayor_count extends OptionalIntField(this)
  object aliases extends MongoListField[Venue, String](this)
  object popularity extends MongoListField[Venue, Int](this)
  object popularityUpdated extends OptionalJodaDateTimeField[Venue](this)

  object tags extends MongoListField[Venue, String](this)
  object categories extends MongoForeignObjectIdList(this, Category)
}

object Venue extends Venue with FSMetaRecord[Venue] {
  override def collectionName = "venues"
  def createRecord = new Venue
  override def mongoIdentifier = NamedMongoIdentifier.venue
}
```

lift-mongo-record & querying III

```
// Foursquare's query engine allows for fluid queries in code
Venue where (_.venueName is "Starbucks")
Venue where (_.venueName nin ("Starbucks", "Whole Foods"))
Venue where (_.geolatlng near (40.72, -73.99))
```

- *Thank you* to @j1iszk4 for sharing this!

Outline

- 1 Introduction
 - What is Scala?
 - Java <-> Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers

- 2 Scala + MongoDB == Win
 - lift-mongo
 - **casbah**
 - STM + MongoDB via Akka

- 3 Closing

Shameless Self Promotion

- Why Casbah?
- Background in pymongo + MongoKit
- Java driver too... “Java-ey”
- Didn't quite “get” scamongo and mongo-scala-driver early on
- scamongo's base didn't fix most of my issues w/ the Java Driver (just helped connection management)
- scamongo's ORM libraries were dependent on Lift (now scamongo is defunct and has become lift-mongo)
- mongo-scala-driver's shapes, etc were *very* confusing to me as a newbie w/o much functional background

Casbah is Born

- Borrowed bits I liked/understood from other places and built something that felt comfortable to me
- Early on, very pythonic
- Query DSL, grown from wanting a feel close to the “metal” based on generic MongoDB knowledge
- Heavily influenced in structure by @jorgeortiz85's libraries
- Quickly grew as I used more and more MongoDB with Scala; features have been grown organically from my own needs.

Interacting with DBObjects I

- `DBObject` is far too structurally Java.
- Sought to make them more usable & readable from Scala
- Most recently - match Scala 2.8 collection Factory/Builders
- Implicit conversions of `Product` (base for `Tuple`), `Map`. Explicit method `asDBObject` for corner cases.
- 'Pimped' version of `DBObject` via `MongoDBObject` - lets `DBObject` implement Scala's `Map` trait.

Interacting with DBObjects II

```
import com.novus.casbah.mongodb.Imports._ // Only import needed - mongoDB type
aliases imported too

val coll = MongoConnection()("test")("testData")

// Map
val map: DBObject = Map(
  "foo" -> "bar",
  "spam" -> "eggs",
  "up" -> "down",
  "pie" -> List(
    "cherry",
    "blueberry",
    "apple",
    "rhubarb",
    "3.14"
  )
)

// 'Product'
val product: DBObject =
( "foo" -> "bar",
  "spam" -> "eggs",
  "up" -> "down",
  "pie" -> List(
    "cherry",
    "blueberry",
    "apple",

```



Interacting with DBObjects III

```
        "rhubarb",
        "3.14"
    )
).asDBObject // Explicit conversion method

// "Factory" method
val constructed: DBObject = MongoDBObject(
    "foo" -> "bar",
    "spam" -> "eggs",
    "up" -> "down",
    "pie" -> List(
        "cherry",
        "blueberry",
        "apple",
        "rhubarb",
        "3.14"
    )
)

// We showed the builder before
val builder = MongoDBObject.newBuilder
builder += "foo" -> "bar"
builder += "spam" -> "eggs"
builder += "up" -> "down"
builder += "pie" -> List("cherry", "blueberry",
                        "apple", "rhubarb", "3.14")

val built: DBObject = builder.result
```

Interacting with DBObjects IV

```
// Also responds to the 'Map' methods...  
built += "x" -> "y"  
built.getOrElse("x", throw new Error("Can't find value for X"))  
/* res15: AnyRef = y */
```

- **DBCollection** behaves as a Scala `Iterable`, but interaction is mostly the same (with addition of methods like `+=`).

Fluid Query Syntax I

- My thought: Instead of keeping track of **Yet Another API**, MongoDB's Query Objects should "just work".
- Two kinds of Query Operators - 'Bareword' and 'Core'.
- Bareword Operators can be started as 'bare' statements:

```
val setMulti = $set ("foo" -> 5, "bar" -> "N", "spam" -> "eggs")
/* setMulti: DBObject = { "$set" : { "foo" : 5 , "bar" : "N" , "spam" : "eggs"}} */
val pushAll = $pushAll ("foo" -> (5, 10, 15, 20, 25, 38, 12, "bar", "spam", 86,
    "eggs", "omg", 412, "ponies"))
/* pushAll: DBObject = { "$pushAll" : { "foo" : [ 5 , 10 , 15 , 20 , 25 , 38 , 12 ,
    "bar" , "spam" , 86 , "eggs" , "omg" , 412 , "ponies"]}} */
```

Fluid Query Syntax II

- Core Operators need to be anchored to the right of a `DBObject` or a `String` (typically representing a field name):

```
// Find any documents where "foo" is between 5 and 15
val findFoo: DBObject = "foo" $gte 5 $lte 15
/* findFoo: DBObject = { "foo" : { "$gte" : 5 , "$lte" : 15}} */
// Find any documents where "bar" contains 1, 8 or 12
val findIn: DBObject = "foo" $in (1, 8, 12)
/* findIn: DBObject = { "foo" : { "$in" : [ 1 , 8 , 12]}} */
```

- Just a small taste - all MongoDB Query Objects are supported (For 1.4.x syntax - 1.6.x (\$or, etc. soon))

Other Features I

- Custom converter implementations which allow most Scala types to be serialized cleanly to MongoDB. (Joda time serialization/deserialization support).
- Improved GridFS Functionality (loan pattern, support for `scala.io.Source`)
- Wrapper objects for Map/Reduce system (Help parse results to warn of errors, etc)

Coming Soon I

- Max Afonov @max4f working on annotation driven object mapping.
- Investigating ActiveRecord implementation, with fluid query syntax support.
- Support for MongoDB 1.6.x features.

A Taste of Casbah's ORM I

```
import scala.reflect.BeanInfo

import com.novus.casbah.mongodb._

import Imports._
import Implicits._
import mapper._
import annotations._

trait Identified {
  @ID(auto = true) var id: ObjectId = _
}

@BeanInfo
class Agency extends Identified {
  @Key("agency_id")      var name:      String = _
  @Key("agency_name")    var description: String = _
  @Key("agency_url")     var url:      Option[String] = None
  @Key("agency_timezone") var tz:      String = _
  @Key("agency_lang")    var lang:     Option[String] = None
  @Key("agency_phone")   var phone:    String = _

  override def toString = "Agency(name = %s, description = %s, url = %s, tz = %s, lang = %s,
    phone = %s)".format(name, description, url, tz, lang, phone)
}
```

A Taste of Casbah's ORM II

```
object Agency extends Mapper[Agency] {  
  conn = MongoConnection()  
  db = "nyct_subway"  
  coll = "agency"  
}  
  
val mta = Agency.findOne(new ObjectId("4c61aecb6f9ee7cdad5b0073"))  
// => Option[Agency] = Some(Agency(name = MTA NYCT, description = MTA New York City Transit,  
  url = Some(http://www.mta.info), tz = America/New_York, lang = Some(en), phone =  
  718-330-1234\n))  
  
val bart = new Agency  
bart.name = "BART"  
bart.tz = "Same as Twitter"  
bart.description = "The subway in SF"  
bart.lang = Some("pig latin")  
  
val bart_as_dbobject = Agency.asDBObject(bart)  
// => com.novus.casbah.mongodb.Imports.DBObject = { "agency_name" : "The subway in SF" ,  
  "agency_timezone" : "Same as Twitter" , "agency_id" : "BART" , "lang" : "pig latin" ,  
  "_id" : { "$oid" : "4c61b568b24ad2b175268dff"}}  
  
val barts_new_id = bart.id  
// => com.novus.casbah.mongodb.Imports.ObjectId = 4c61b568b24ad2b175268dff  
  
val bart_saved = Agency.upsert(bart)  
// => Agency = Agency(name = BART, description = The subway in SF, url = null, tz = Same as  
  Twitter, lang = Some(pig latin), phone = null)
```

A Taste of Casbah's ORM III

```

val bart_reloaded = Agency.findOne(new ObjectId("4c61b4bdb24ad2b172268dff"))
// => Option[Agency] = Some(Agency(name = BART, description = The subway in SF, url = null,
    tz = Same as Twitter, lang = Some(null), phone = null))

@BeanInfo
class Route extends Identified {
  @Key("route_id")          var name:      String = _
  @Key /* infers key from field name */ var agency_id: String = _
  @Key("route_short_name")  var short_name: String = _
  @Key("route_long_name")   var long_name: String = _
  @Key("route_desc")        var description: String = _
  @Key                      var route_type: Int = _

  override def toString = "Agency(%s -> %s)".format(short_name, long_name)

  // foreign key, anyone?
  lazy val agency = MongoConnection()("nyct_subway").mapped[Agency].findOne("agency_id" ->
    agency_id).get
}

object Route extends Mapper[Route] {
  conn = MongoConnection()
  db = "nyct_subway"
  coll = "routes"
}

//val N_train = Route.findOne(new ObjectId("4c61aecb6f9ee7cdad5b0275"))
//val of_course_its_mta = N_train.get.agency

```

A Taste of Casbah's ORM IV

```
// EVEN MOAR! nested, optional documents? collections of nested documents?
```

```
@BeanInfo
```

```
class Address {
```

```
  @Key var street: String = _ // required strings
```

```
  // optional strings and nulls are stripped from final output
```

```
  @Key var street2: Option[String] = _
```

```
  @Key var city: String = _
```

```
  @Key var state: String = _
```

```
  @Key var zip: Int = _
```

```
}
```

```
@BeanInfo
```

```
class Person {
```

```
  // "_id" can be anything, not just ObjectId
```

```
  @ID var unix_name: String = _
```

```
  @Key var first_name: String = _
```

```
  @Key var last_name: String = _
```

```
  @Key var address: Address = _
```

```
  // optional address. not everyone has a job!
```

```
  @Key var work_address: Option[Address] = None
```

```
  // more addresses, a whole list, empty by default
```

```
  @Key var other_addresses: List[Address] = Nil
```

A Taste of Casbah's ORM V

```
}

val home = new Address
home.street = "1 Main Street"
home.city = "Brooklyn"
home.state = "New York"
home.zip = 11201

val work = new Address

val joe_sixpack = new Person
joe_sixpack.unix_name = "jsixpack"
joe_sixpack.first_name = "Joe"
joe_sixpack.last_name = "Six Pack"
joe_sixpack.address = home

joe_sixpack.work_address = Some(new Address).map {
  work =>

    work.street = "25 Wall Street"
    work.city = "New York"
    work.state = "New York"
    work.zip = 10001

    work
}

joe_sixpack.other_addresses = home :: work :: Nil
```

A Taste of Casbah's ORM VI

```
object Person extends Mapper[Person]
Person.asDBObject(joe_sixpack)

/*
{
  "unix_name" : "jsixpack" ,
  "first_name" : "Joe" , "last_name" : "Six Pack" ,
  "address" : {
    "street" : "1 Main Street" ,
    "city" : "Brooklyn" ,
    "state" : "New York" ,
    "zip" : 11201
  },
  "work_address" : {
    "street" : "25 Wall Street" ,
    "city" : "New York" ,
    "state" : "New York" ,
    "zip" : 10001
  },
  "other_addresses" : [
    {
      "street" : "1 Main Street" ,
      "city" : "Brooklyn" ,
      "state" : "New York" ,
      "zip" : 11201
    },
    {
      "street" : "25 Wall Street" ,
      "city" : "New York" ,
```

A Taste of Casbah's ORM VII

```
    "state" : "New York" ,  
    "zip"   : 10001  
  }  
]  
}*/
```

Outline

- 1 Introduction
 - What is Scala?
 - Java <-> Scala Basics
 - Implicits and Pimp Hats
 - What is MongoDB?
 - A Taste of MongoDB
 - MongoDB + Scala Drivers

- 2 **Scala + MongoDB == Win**
 - lift-mongo
 - casbah
 - **STM + MongoDB via Akka**

- 3 Closing

STM + MongoDB via Akka I

- Akka has an implementation of STM inspired by Clojure's; allows datastructures such as Maps and Vectors to become transactional.
- Akka STM supports persistence to several backends including MongoDB.
- Allows you to setup relatively simple, code managed concurrent transactions with state stored safely in MongoDB.
- Supports JTA; not yet distributed (Dependent on Multiverse, which is working on distributed STM)

Links

- **mongo-scala-driver** <http://github.com/alaz/mongo-scala-driver>
- **lift-mongo** <http://www.assembla.com/wiki/show/liftweb/MongoDB>
- **FourSquare's Lift Mongo DSL Code . . . coming soon?** @jliszka
- **Casbah** <http://novus.github.com/docs/casbah>
- **Jorge Ortiz' (@jorgeortiz85) Libraries**
 - **scala-javautils (Scala 2.7.x)** <http://github.com/jorgeortiz85/scala-javautils>
 - **scalaj-collection (Scala 2.8.x)** <http://github.com/scalaj/scalaj-collection>
- **Recommended books...**
 - **Programming Scala (Subramaniam, Pragmatic Bookshelf, 2009)**
 - **Programming Scala (Payne & Wampler, O'Reilly 2009)**

Contact Info

- Twitter: @rit
- Email: bwmcadams@gmail.com
- Github: <http://github.com/bwmcadams> | <http://github.com/novus>
- IRC - [#mongodb](http://freenode.net)
- MongoDB Mailing List <http://groups.google.com/group/mongodb-user>
- Casbah Mailing List <http://groups.google.com/group/mongodb-casbah-user>
- Boston MongoDB Conference - Sept. 20 (Cambridge, Mass.)
<http://10gen.com/conferences/mongoboston2010>
- MongoDB NY Users Group
<http://www.meetup.com/New-York-MongoDB-User-Group/>